

AI Mania Is Eviscerating Global Decisionmaking

By Nikhil Suresh

AI Mania Is Eviscerating Global Decision-Making

on 18/07/26

I strongly believe there are entire companies right now under heavy AI psychosis and it's impossible to have rational conversations with them about it. I can't name any specific people because they include personal friends I deeply respect, but I worry about how this plays out.

Mitchell Hashimoto, of HashiCorp and Ghostty fame

Over the past year, I've run point on all of our company's sales, led the technical components of all but two of our engagements, and over the lifetime of this blog have had something like 300 catchups with professionals from around the world. This has ranged from people on the ground in niche service industries to executives at Fortune 500 companies.¹ Also, and this is 100% true, Matt Mullenweg once asked me for coffee because he read the [AI piledrive essay](#), and in context probably enjoyed it, but had to cancel because he hadn't realized he had a flight later the same day. I am willing to pay a competent witch to hex him for this slight. Because of this, I've had a front-row view to our collective institutions across both the private and public sector undergoing breath-taking mass psychosis. This essay is an attempt to describe the bizarre dynamics that are currently at play, as I am in the rare position where my wellbeing is not contingent on paying lip service to madness, and to reassure the people trying to survive amidst all of this that they are not crazy.

The reality is thus: the people in charge either have no plan, or see no path forwards other than keeping their heads down. Not at banks, not at hospitals, not in our government institutions. The world's organisations have been captured by people in the throes of frothing excitement, and saner people who now live in a state of constant commingled fear and frustration.

I. AI Investments Are Generally Total Failures

Reading this while working for a division that pivoted to provide interfaces for agentic workflows, only to discover that only ten users had ever touched the products we made for agents, only to pivot again to support for agentic workflows, which has a lot of competition because every company has to do something agentic now and there's only like four things you can do in that space, is bracing.

An editor of this essay

Are companies actually seeing massive productivity gains from their AI adoption? Does any of this sordid affair *make sense*?

This should be an easy question, but it is surprisingly hard to get a straight answer to it. Executives

that tell the press that their company has gone insane will quickly find themselves removed from their positions. Employees who are honest will find themselves fired in short-order, or “randomly” selected for a round of layoffs. In fact, it is in the interests of almost every actor in the space – boards, executives, employees, vendors, consultants – to obfuscate and misrepresent the success rate of AI projects. Many publicly traded companies are putting out announcements about their AI productivity gains when I know for a fact that the businesses have done nothing other than purchase Copilot licenses and declare victory.

Yet we need to know if these projects are panning out – if the total focus on AI as a core tenet of business strategy is succeeding at a reasonable rate, then a discussion about the relative risk and reward is warranted.

Unfortunately, we live in a dark timeline. All of the AI projects we have observed as a team are failing. Every single one – we have seen 0% success in a year and a half, not only amongst projects we have been asked to participate in²² We have rejected all AI implementation work. It is absolutely a gigantic bubble and we have minimized our exposure to it – every single one of our current contracts would be totally unaffected by OpenAI collapsing, save for *perhaps* some second-order effects such a recession causing a client to become unable to pay us. And there's nothing we can do to insulate ourselves from that anyway., but even within projects that we have observed in passing while doing totally unrelated work. Even if you grant that AI tooling accelerates specific workloads, the method and scale of the current investments is senseless. Frequently the failure is not related to AI itself, but rather that companies are terminally bad at running software projects effectively, and as [I have remarked previously](#), AI projects are subject to all the failure modes of normal projects *plus* you can get everything right and then still fail because of the method's novelty. Very few companies are so good at shipping software that they can afford the extra risk profile.

Often enough, though, it's an actual failure in what LLMs can accomplish. The most common version of this, being rolled out across businesses around the world, is the internally-facing chatbot, or for the more daring company, the customer-facing chatbot. The story is always the same. For the former, I've never seen substantial internal uptake from inside a business. Employees don't use internal chatbots because companies tend to have low-quality documentation and an LLM is not psychic – it can only know things that have been written down and made accessible. For the latter customer-facing applications, I have rarely had a pleasant experience as a consumer, with *perhaps* the exception of live transcription during medical appointments – hardly something worth pivoting an entire organisation around. In both cases, project leaders are very careful to avoid tracking basic metrics, such as whether the tools are being used at all, or they track metrics that are easily gamed.

For example, my last consumer interaction was attempting to get help from Mitsubishi following an automotive failure, where a very polite robot asked me to describe the problem and that I'd receive a call back as soon as someone was available. This was the single most competent implementation of such a project I've seen in the wild, in that the voice was natural sounding, responded quickly, was clearly “live” in production, and promised a swift resolution.

That was six months ago, and I did not, in fact, get a call back.

When Mitsubishi did not call me back, what happened? Did that request just go into the void, showing one less incident for the year? Does it appear that the phone bot resolved my query without the need for human intervention? All we know is that it didn't show up as an error, or I'd have received a call. I'm sure it looks great in all sorts of ways except the one that matters, which is that I was planning to buy a car and decided *not* to buy another one of theirs.

For this reason, our team has quickly learned while on an engagement not to ask anything about ongoing AI projects in any context – by the time that project has started, it is too late for the management team, and intervention is not possible until a crisis point is inevitably reached. There is no conceivable positive outcome. The failure rate is so high that even basic inquiry leaves us in an untenable position. Any coherent question about how it's going, what the goal is, who is using it, constitutes an inadvertent attack on the chain of command responsible for the work because *there are no good answers to anything*. Even in rare cases where my interlocutor has stated that things are going well (usually while the project is still mid-flight and failure has not had a chance to manifest), it is generally obvious that they are doomed, but at least in these cases I can simply agree and then go home to scream into a pillow for six hours straight.³³ One of the most valuable rules I've heard, from Gerry Weinberg, is that consulting is influencing people *at their request*. Unless someone has indicated that they want us to stick my nose in, usually by explicitly saying they want guidance on general data strategy, we just let the projects fail in peace. You can barely recognize me, I'm so calm these days.

All of this is to say that I am very confident that almost every report at a company about “massive AI productivity gains” is untrue as a matter of brute fact. Even if some companies are seeing clear gains, this is the exception, not the norm. With that assumption in place, we can talk about the dynamics at play, and how it has become impossible for many organisations to stay focused on things that actually matter to their long-term (or even short-term) health.

II. Heretics Will Be Shot

It has become outright dangerous to even raise the possibility that AI might not be the solution to a problem, let alone be the sole focus of a company's entire strategy.

In every sufficiently large business we have observed (say, with 500+ employees), we have noted that continued advancement, and increasingly continued employment, has started to require repeated professions of belief in the transformative power of AI for said business. I am not talking about providing ideas about how to use AI in the business – I mean *religious* profession, declarations of faith. Overwhelmingly these statements are made by non-technicians, though it is not uncommon for technicians to emit deranged statements to curry favour.

There have been several occasions where I have seen someone, apropos of nothing, blurt out almost word-for-word “AI is changing everything”, only to concede moments later that their organisation does not currently use LLMs for anything, and indeed, that they cannot name a single thing that has changed other than they get some use out of ChatGPT (frequently the free-tier). In one extreme case, I have seen an executive confess that they had never even used ChatGPT or any AI tool in their life, immediately after producing a technical strategy for an organisation with \$2B+ in revenue which was entirely centered around AI.

Initially these statements were so absurd on their face that I thought it was some cynical ploy to achieve thought leader status, and there are certainly some people doing this – I have had it admitted to me. But the broader reality is so much worse: people who have no background in the technology at all *actually believe what they are saying*. As a general rule you should avoid getting into business with a liar, but if you *must*, you can at least reason with them even if only in private. A true believer is much more threatening because they are impervious to even inducement by self-interest.

The turning point in my belief was watching someone with a spectacular amount of money on the line fire their highest performers because they were achieving that performance without LLMs. When an employer *publicly* talks about AI innovation, we have to ask ourselves if they're simply trying to manipulate the market or customers. When they *privately* commit to strategies like this with their own money at stake, with no attempt to communicate that strategy to external clients, I can only assume they really mean what they're saying.

A while ago, I wrote [“Contra Ptacek’s Terrible Article On AI”](#), which was focused on the fact that many of Ptacek’s points in his own essay [“My AI Skeptic Friends Are All Nuts”](#) were internally inconsistent.⁴⁴ We have since kissed and made up in private, though I don't think we've budged at all on the core points of our viewpoints. I maintain that Thomas is [a very talented writer with a lot of good advice](#) who just happened to blow it massively that one time because he takes Hackernews commenters too seriously. We all have our weaknesses. Mine is people telling me that "Scrum is good if you do it right". But on the crux of the matter, we are actually in total agreement, because he opens his essay with this:

Tech execs are mandating LLM adoption. That’s bad strategy.

Which is to say that we can sidestep arguments about the precise utility of LLMs entirely and we're left in a very simple place – it is entirely obvious to both myself and Ptacek, two people that are coming at this from fairly opposed views, that people are being really, really stupid about this, *and* that organisations are demanding bizarre workflow constraints from their specialist staff.⁵⁵ This is always baffling to me as a matter of being a responsible adult. If I was somehow CEO at a hospital or civil engineering firm, I would not for a second think it's my place to start mandating specific procedures or building techniques without explicit agreement from the professionals on staff – how fucking clueless are the non-technicians who have attended a few talks and are now making mandates about how their extremely expensive professionals are doing their jobs?

These mandates have led to extremely strange places. Several of my peers now “AI-wash” their work, meaning that even when they can perfectly competently execute on their jobs to the satisfaction of their management teams, said managers are unhappy if the engineers haven't used AI in the work... so now they're lying about using LLMs even in contexts where their professional judgement is that they aren't the appropriate tool. They just do the work, the same way they have for decades, and say Claude did it. Others are being measured on their AI bills with “token leaderboards”, where higher is *better* because I have evidently fallen into the pocket of Hell where the demons torment me by doing elaborate impressions of absolute fucking morons, so the people hired for their freakish ability to perform system optimisation do the obvious thing. They set the LLMs prompting themselves in a semi-plausible loop in case someone inspects the token consumption and then they watch Netflix. Not a single one has been caught, even when their own assessment of the output is that it isn't suitable for deployment.

Checking out a parallel copy of our Go repository and telling the AI to rewrite the whole thing in Zig while I work on something else just so I can keep my job. I hate this shit so much. My job has usage tracking and quotas. I don't use it for actual work, I just spin it up and disregard the output.

An actual software engineer

In fact, the only people I know of to be fired over this whole thing are people that have expressed visible doubt about this organisational strategy, which again, even Ptacek thinks is transparently

dumb. The net result is that everyone has learned very quickly to praise executives on their visionary AI prowess, or they will be gunned down in the proverbial streets.

III. AI Demos Are The Mind-Killer

Bless me, Father, for I have sinned. It has been ∞ days since my last confession. I accuse myself of the following sins:

One of the main pieces of infrastructure we deploy at our clients is an analytics-focused database called Snowflake – for a typical business, the bill is *tiny* because it's a pay-as-you-go situation and we can process all their data in one minute a day, you get a very hands-off deployment, and in short it has many characteristics that are very pleasant for our work. One of the features in Snowflake that we *don't* use is called Cortex.

Cortex is their AI chatbot layer, with the ability to plug into metadata (for non-nerds, descriptions of your data, like what a column in a spreadsheet means) and query a company's database autonomously. In theory, you can ask a question like "What was our revenue for last week?" and it will spit out an answer.

It is not really suitable for production usage. From memory, the last time I was given a presentation on it, by actual Snowflake staff, they reported that ideal configuration results in something like ~92% accuracy due to the complexity of data at a large business (see: probably best-in-class for these tools, but imagine your CFO having one in every ten of their numbers be outright wrong) and there were serious issues with managing deployments. Nonetheless, it can be used to produce some very flashy demonstrations.

On several occasions, we've been exposed to folks that have been sort of lukewarm on our main offerings, but they *really, really* wanted to use AI to perform a natural language query on their data. And we thought "Okay, if you *really* want to see it, maybe we can caveat this appropriately and show you what it might look like."

This was a terrible mistake. It backfired in the most predictable way imaginable – every lukewarm client that saw the chatbot in action, *even with us telling them that it was not going to accomplish what they wanted*, wanted to buy it immediately. Every other consideration, including millions of dollars that we could plausibly help them achieve by non-AI means, was swept aside. It was like a dark and terrible force seized control of their limbs, plunged their hands into their own chests, and presented their still-beating credit cards to us in grim supplication. We were so mortified by the inexplicable shift in energy that we (wisely) declined to take the money and ended the sales process, and soon thereafter removed Cortex from our list of demonstrations. It would have been too irresponsible to exploit this gap in their reasoning, and frankly, it was already irresponsible to have even run the demonstration – doctors don't walk around showing off cool pills that they'd never prescribe.

Watching the total 180°, that shift from ice-cold to red-hot buying frenzy, was a deeply unsettling experience. It was personally uncomfortable to see people that clearly didn't gel with us interpersonally suddenly dying to enter an ongoing relationship, but more broadly uncomfortable because for a brief moment I began to understand what is happening in sales meetings around the world. There was no warning I could have given that would have made them refuse to buy the damn thing – their appetite was as large as their budget could stretch, and some part of me wonders if this is because they knew that their ravenous hunger would be present in their own customers.

They'd just buy it from us, then pivot right to a larger company and mind control *their* leadership team until the buck finally stops with the loser that needs to justify the expense. The main protection against this seems to be that the median vendor is [so bad at their jobs](#) that we had presented the first even somewhat-working products these people had seen, and this *included an ASX-listed company that was already bragging about their AI usage*. It took our team two hours to produce something that was frankly not that good – basically just typing text descriptions of data into a web browser – and it was still better than anything the leads had seen because they *had nothing* to show for all the investment.

In fact, we have been forced to opt out of every sale where the lead has expressed anything beyond the most fleeting curiosity in the use of AI in their business. I don't mean that we've heard that they're interested in AI and elected to drop the contract on moral grounds. I mean that, over the course of the engagement, these people have exhibited a pattern of behavior that has made it near-impossible to sell to them without incurring reputational and legal risk, and are furthermore crafting management environments that I can only describe as cultish, ineffective, and “please dear God, do not let it be on earth as it is on LinkedIn”.

IV. Executives, Game Theory, and The Emperor's Clothes

The good news is, CISOs are used to having to protect the business from their hare-brained initiatives, and this one isn't really that different, except that there's a cult-like atmosphere to it that you didn't see with, say, the cloud. It almost doesn't matter whether you embrace the initiative or not; there's work to be done to manage the risk, so that's what you do. From talking to CISOs everywhere, I would say most of them are quietly skeptical but afraid to speak up.

Career CISO and well-known speaker that asked to remain anonymous

Despite the substantial prevalence of true believers, many of the people running large AI initiatives, or making public statements about them, do not believe what they are saying. There are “heads of AI” who read this blog, at companies with \$1B+ in annually recurring revenue, who have written in to say they believe their job is totally fraudulent but it was the only promotion pathway remaining at the organisation.

On a trip overseas, I had the privilege of a meeting with one of the Fortune 500 executives mentioned at the beginning of the post, who will remain anonymous so that they are not executed by firing squad by their board. As we were chatting, it became clear that they were very switched-on and technically competent, and they *also* happened to be at a company that had committed to the usual battery of exorbitant claims about their recent innovations – we've 100x'd our productivity, AI is the future of everything, I am but a vessel for OpenAI to make love to my wife. You know, normal things. But since I had them there without any microphones around, I asked *why* this was being repeated without opposition. Was it just sales fluff?

The answer was a lot more interesting. It was *partially* ridiculous sales material being delivered to an easily excitable audience, but this was not the dominant factor constraining honesty. Executives at their *customers* were saying absurd things about achieving 100x productivity, and this meant that if any executive at the *vendor* said that these gains were not plausible, it would undermine the credibility of the customer's executive, be perceived as an attack (or heresy), and possibly result in an enterprise contract cancellation. And getting enterprise contracts cancelled because you wanted to opine on something that doesn't really matter to your organisation's mission is a great way to get

fired.

But this company was also a major player, of the kind that signs enormous enterprise contracts with other companies. So presumably there is *another* vendor that has sold to them, and *their* CEO is worried that saying something sane will contradict *this* executive, and very quickly we can see how we can have executives around the world nervously pointing guns at each other, not wanting to be shot first but also watching everything gradually spiral out of control.⁶⁶ If you're an executive, board member, or anyone in charge of an "AI project" that feels trapped, I would love to hear from you. I will file the serial numbers off any stories very carefully, as I've done here and in every other article. This is to say that we're facing a [coordination problem](#) around executives being honest around the AI gains they've witnessed – if they co-operate, they keep their jobs. If they defect, they will possibly be fired by their embarrassed peers (who have now been implicitly called liars, cowards, or incompetents) and then replaced with someone that will toe the line anyway. If they could *all* admit the truth at once there might be some hope, but there is no way to coordinate that event.

This sounds deeply concerning, but it is worth noting that it means that some executives who are emitting nonsensical statements are not as dull as they might seem at first – they're in a fraught political environment, where they are surrounded by many people that are gunning for their roles, and subject to the whims of a board that is undergoing similar pressure. Against all the dictates of reason, I have presented on navigating AI hype to people on S&P 500 boards⁷⁷ This sounds very fancy, but I think it was secretly one of those compulsory professional development things and half the audience were just like, making dinner. Truly, HR and professional bodies make victims of us all. and they are in exactly the same situation – the main comments I remember from the session were board members admitting they were skeptical, but expressing anxiety that their positions were contingent on demanding AI investment. One of them commented “investing this early seems like risk without much upside”. About two years later, I can see now that their decade-old multi-billion dollar organisation is now branded as “AI-native”, whatever the hell that means.

V. You Must Be This AI-Native To Ride

All of the above converges on the state that we find ourselves in now, where effective decisionmaking has ground to a halt. Collectively, what started as a few people undergoing either destabilising psychological events or being caught up in hype has now resulted in an environment where leaders cannot speak honestly about their beliefs on how best to guide organisations, for fear of being removed, creating a sort of distributed government by assassination. This means that the least sensible recommendations are going totally unchallenged, resulting in employees being evaluated on totally gameable metrics such as “money spent on AI”, and those employees must play along to avoid being terminated. This has also created an insatiable appetite for purchasing “AI” solutions, which target both true believers that will believe implausible claims, and also non-believers that cannot decline the purchases without having their commitment to the cause coming into question.

This means that all offers that are subject to internal politics at an ideologically captured organisation must include AI alignment, even if the value proposition is patently ambiguous. My assessment of the market so far is that a substantial component of the outburst of AI projects are actually non-AI projects with an AI element slapped on after the fact to pass the purity test.

For example, I recently witnessed an organisation handling a database migration from an Oracle

database to Snowflake – instead of handling the migration directly, the vendor bolted on a preliminary phase which involved trying to get an LLM to automate the translation of the Oracle-flavored SQL to Snowflake-flavored SQL. When the project failed (due to issues getting enough permissions to automate the work, not because an LLM can't do something that easy), the vendor simply started handling the translation by hand but the company billed it as an AI-driven success because some inconsequential portion of the SQL had been translated by AI before being pasted over.

What was actually purchased? A totally standard database migration to help an executive meet the strategic deliverable of decommissioning a system prior to license renewal. What was sold to their superiors? “I allocated a substantial percentage of my budget to AI and it helped me accomplish my mandate.” True AI projects, of the kind that is driven by an LLM as the sole mechanism underlying it, where the project can clearly fail to deliver specific numbers, are actually very rare. We mostly see them in the context of startups, and frankly we have stopped engaging with them because we kept getting to the end of the sales conversation and finding out *they wanted us to build the product that they were marketing as completed*.

However, some projects simply do not have an easy way to tack on the AI label, or the person advocating for them either does not want to lie or has not understood that lying has become necessary. In all cases, this either kills the request for funding outright, or adds a pervasive and intractable drag on all communications, as every request must be worked and re-worked until it is “AI enough”. Failure to comply will either result in denial or, in many cases, a demand from a true believer to know why the extra work “can't be done with AI”. Many companies have actively publicized that this is their new hiring policy – when a member of staff requests additional headcount, they must demonstrate that they have tried to use AI first. The part that's being left out is that if you say you used AI and still need the help, you will be labelled “bad at AI” and potentially laid off.

The net result of this is that almost every large organisation that I am aware of is no longer able to focus on anything important, unless they are one of the (very) few organisations where AI happens to address their highest priorities. They cannot buy sensible software, hire competent talent, communicate honestly with executives about the state of projects, or undertake any sort of sensible initiative.

VI. Navigating AI Mania

An emptiness falls through you
As you realize what this means
You're starting to feel what I feel
Now you've seen what I've seen

So Sick, Domesticated Incels

This is an unfortunate situation to be in, but it will pass eventually. I've learned a lot about the latent insanity that we have inculcated in our leadership strata, and unfortunately those traits will persist long past the current bubble, merely awaiting another similar reactivation trigger – and some organisations will stay captured until they have totally collapsed, in the way that not everyone has successfully moved away from the dreadful blockchain affair. That's something to write about for another time.

What I wanted to get to were some thoughts on surviving the immediate crisis, either by directly making systemic improvements or by holding onto your sanity. I'll start with the "making improvements" part, because that's the situation I find myself in the most frequently.

When You Have Another Objective

We're going to do a lot of sucking it up and smiling here. This section assumes that you are trying to achieve some goal that *isn't* repairing the organisation's manic stance, but either trying to course-correct a specific project (and possibly risk getting fired as either a leader or consultant) or achieve some totally unrelated goal.

1. Where possible, when raising issues, do not have conversations about the state of AI projects in group settings, as this creates a dynamic where each individual member of the group is worried about outing themselves in front of their peers. Arrange for one-on-one settings. Make it clear that you are willing to countenance that the current AI environment is frothy, and that you will keep opinions unidentifiable when raising them elsewhere. **Be extremely aware that the most outspoken people can be identified by their peers, so take care to avoid exposing your sources by, e.g. direct quotes.** In the event that only a small minority (say, one person in a group of six people) is willing to speak out, it might be worth giving up and moving on to a patient that has better chances.
2. For ongoing projects, an effective trick that I believe I picked up from *Secrets of Consulting* is the anonymous poll, where you can ask individuals to rate their opinion of an AI project's success chances on a scale of 1 to 10. The typical split I have observed is half of those involved rating the project at a 3/10 and others at around an 8/10 – a clear bimodal split *on a project that was already three years late*. Bringing this data to a CEO can be an effective method of pointing out that some information is clearly being hidden from them on the state of the project.
3. Always involve people on the ground. The only source of data on whether projects are succeeding or the investment is going anywhere are the people that use it for their day-to-day activity. Care must be taken to bring them into the environment where they are treated with respect (all sufficiently large companies have people that view subordinates as not-quite-real-people). It is not uncommon to uncover worldview-shaking information in short order – with one client, we uncovered that staff were *totally unaware* they had been given licenses for AI tooling, which cast into doubt all productivity claims.
4. *Do not* question the broadest claims about AI. I cannot emphasize this enough. If someone says "AI is changing everything", just let it pass if your goal is to fix an object-level problem rather than challenge the reality at the institution. The challenge can only come after you have gained the trust of the most senior person involved. Trust is gained over a meal in private where you assuage their anxieties, not by embarrassing them in front of peers.
5. Remember that you do not know what statements have been emitted prior to entering a room. There will sometimes be people that have publicly committed to statements like "I am 100x more productive than I was last year", and some may even wish they hadn't said that but are too embarrassed to walk it back. In an untested room, common sense like "LLMs should not be allowed to deploy code without human review" can kill your chances to make an impact before you've even started.
6. My practice requires me to maintain an honest relationship with my clients or the whole thing falls apart, so I can't do this – but honestly, if you work in the fire service and need money to stop a puppy from catching fire, just lie. It's fine. History will forgive you. Add a \$10,000 AI chatbot to your project, exclusively discuss that part in meetings, *whatever*. Save that puppy.

When You're Just Trying To Survive

This is for people that are just waiting for the bubble to burst and trying not to go nuts.

1. I have bad news – accept that you are probably not going to meaningfully push back on any of this. This is not a feature of AI, it's a feature of dysfunctional companies.
2. If you feel like you're going absolutely nuts, consider switching over to contracting. I've advocated for contracting many times over full-time employment, but you'll get paid a lot more and be left out of most internal politics. Also when you run into a really intolerable situation, you'll know that you've got a fixed end-date.
3. I do my best to limit my uptake of AI-related news, as it is pretty crazy-making and unproductive to consume. I no longer visit Hackernews, Reddit, or really anywhere where I am going to be drip-fed nonsense, though I allow myself exceptions for very funny things like [Apple suing OpenAI over alleged corporate espionage](#). Consume exactly the amount you need to feel like you aren't going insane, then *stop*. Ditto for complaining with friends – and tell them that's why you're talking about it, which buys a lot of tolerance.
4. When someone tells me they are using AI for something when they really shouldn't be, I smile and nod as long as they are unlikely to get themselves killed. Even family. *Especially* family.
5. When someone asks me for my opinion of AI as a programmer, I recommend saying “Oh, that stuff is pretty overblown” and then changing the topic, unless they are in a position where their opinion might influence something important. Non-programmers need this guidance the most.
6. If you're being asked to review huge volumes of terrible AI code, just assume that the organisation is going to burn you out and fire you. You will not convince the person drowning you in 2000 line PRs to stop. Start looking for a new job as if you have already been fired. I have seen this happen many times now, and it always plays out the same way – do the job search while you have energy. Don't worry if your speed drops or management gets annoyed at you. There is no way to avoid that, you can simply choose whether it happens now because of your job search, or later because you are too depressed to work anymore.
7. If your manager is responding to you with clearly AI-generated text, use AI to respond to save your sanity and then look for a new job. Many people assume they will get in trouble for being that obviously rude. You will not, this particular behavior is exhibited only by true believers, and they actually *like* that you've clearly not bothered to engage with them. I know, it's fucking wild.
8. If you're being asked to max out on token usage, look for a new j – okay look, you get it, right? Go find a job that isn't going to wrench reality from your tenuous grasp. They do exist, largely at companies so small that they don't turn up on job platforms. It might take months to find one, so start now.

Fight the good fight, and don't let the bastards grind you down. Godspeed.

1. Also, and this is 100% true, Matt Mullenweg once asked me for coffee because he read the [AI piledrive essay](#), and in context probably enjoyed it, but had to cancel because he hadn't realized he had a flight later the same day. I am willing to pay a competent witch to hex him for this slight. ↩
2. We have rejected all AI implementation work. It is absolutely a gigantic bubble and we have minimized our exposure to it – every single one of our current contracts would be totally unaffected by OpenAI collapsing, save for *perhaps* some second-order effects such as a recession causing a client to become unable to pay us. And there's nothing we can do to insulate

ourselves from that anyway.↵

3. One of the most valuable rules I've heard, from Gerry Weinberg, is that consulting is influencing people *at their request*. Unless someone has indicated that they want us to stick my nose in, usually by explicitly saying they want guidance on general data strategy, we just let the projects fail in peace. You can barely recognize me, I'm so calm these days.↵
4. We have since kissed and made up in private, though I don't think we've budged at all on the core points of our viewpoints. I maintain that Thomas is [a very talented writer with a lot of good advice](#) who just happened to blow it massively that one time because he takes Hackernews commenters too seriously. We all have our weaknesses. Mine is people telling me that "Scrum is good if you do it right".↵
5. This is always baffling to me as a matter of being a responsible adult. If I was somehow CEO at a hospital or civil engineering firm, I would not for a second think it's my place to start mandating specific procedures or building techniques without explicit agreement from the professionals on staff – how fucking clueless are the non-technicians who have attended a few talks and are now making mandates about how their extremely expensive professionals are doing their jobs?↵
6. If you're an executive, board member, or anyone in charge of an "AI project" that feels trapped, I would love to hear from you. I will file the serial numbers off any stories very carefully, as I've done here and in every other article.↵
7. This sounds very fancy, but I think it was secretly one of those compulsory professional development things and half the audience were just like, making dinner. Truly, HR and professional bodies make victims of us all.↵

If you liked this essay, subscribe to our blog to get notified when new essays come out. We're focused on writing helpful, entertaining, and *human* writing, and we simply trust that the marketing will handle itself. We won't flood your inbox with trash. We will never, ever share your email with a third party.